


Qwen3.8 Flash Next : prefix caching works

Single GB10 · 262,144 context · 8 GiB BF16 KV · MTP2 · thinking medium



Measured cold/warm requests; exact answers checked. Cache speeds prefill, not decode.

2026-09-05 · official upstream PR backports, still unmerged · retention every 1,600 tokens